A. Giovanidis 2022

06. Mode-Selection / Cross-Validation

Network Data Analysis - NDA'22
Anastasios Giovanidis

Sorbonne-LIP6

October 25, 2022

Bibliography

A. Giovanidis 2022

B.1 Gareth James, Daniela Witten, Trevor Hastie, Robert Tibshirani.
"An introduction to statistical learning: with applications in R".
Springer Texts in Statistics. ISBN 978-1-4614-7137-0
Chapter 2, Chapter 5
DOI 10.1007/978-1-4614-7138-7

Recap

In the previous course (Regression) we assumed that the real-world model is sufficiently described by a linear model with additive noise:

$$y = \beta_1 x + \beta_0 + \epsilon = f(x) + \epsilon$$

We estimated the unknown β's by the parameters $\hat{\beta}_1, \hat{\beta}_0$.

The following formula predicts for any x

$$\hat{y} = \hat{\beta}_1 x + \hat{\beta}_0 = \hat{f}(x).$$

☞ But we do not know anything about the noise !

Prediction Errors

These predictions cannot be accurate but will always have an irreducible error, no matter how good the choice of the predictor $\hat{f}$:

$$
\begin{aligned}
\mathbb{E}\left[(y - \hat{y})^2\right] &= \mathbb{E}\left[(f(x) - \hat{f}(x))^2\right] + \mathbb{E}\left[\epsilon^2\right] \\
&+ 2\mathbb{E}\left[\epsilon(f(x) - \hat{f}(x))\right]. \\
&= \mathbb{E}\left[(f(x) - \hat{f}(x))^2\right] + \textit{Var}(\epsilon).
\end{aligned}
$$

Two types of errors...

Reducible vs Irreducible

A. Giovanidis 2022

- ▶ The irreducible error, due to the random error ϵ in the model, whose variance is unknown.

- ▶ The reducible error, due to errors in the estimate of the model parameters $\hat{\beta}_i$. This type of error can be reduced by using (a) larger sample sets when estimating the coefficients, or (b) different models $\hat{f}(x)$ that better describe the unknown function $f(x)$ (could be non-linear).

In practice larger intervals are used for prediction to account for both types of errors.

Accuracy

The MSE is a measure of model accuracy.

For the available data set $D_n = ((x_1, y_1), (x_2, y_2), \ldots, (x_n, y_n))$,

$$MSE(D_n; \hat{\beta}_1, \hat{\beta}_0) = \frac{1}{n} \sum_{i=1}^{n} \left(y_i - \hat{\beta}_1 x_i - \hat{\beta}_0 \right)^2. \qquad (1)$$

We guarantee maximum accuracy by

$$\min_{\hat{\beta}_1, \ \hat{\beta}_0} MSE(D_n; \hat{\beta}_1, \hat{\beta}_0).$$

By choosing $\hat{\beta}_1$, $\hat{\beta}_0$ that minimize MSE we reduce the reducible part, but cannot change the irreducible part due to noise.

Train MSE

☞ The predicted response $\hat{y}(x_i) = \hat{\beta}_1 x_i + \hat{\beta}_0$ will be close to y_i, because the parameters are chosen to minimise their difference!
We say that the model is trained with data D_n.

$$MSE_{train} \quad = \quad MSE(D_n; \hat{\beta}_1, \hat{\beta}_0).$$

But! We actually want that the model predicts good values for unknown data, $x_o \notin D_n$.

$$\hat{y}_o \quad = \quad \hat{\beta}_1 x_o + \hat{\beta}_0.$$

Test MSE

A. Giovanidis 2022

We need a different test data set D_m^{test} with a number $m \geq 1$ of samples, to test the accuracy of our prediction model.
For this test data set, we relate the accuracy metric

$$MSE_{test} \quad := \quad MSE(D_m^{test}; \hat{\beta}_1, \hat{\beta}_0) \neq MSE_{train}.$$

Question 1: How good does our "minimum MSE linear predictor" behave for the test data set?

Question 2: If we use other prediction models $\hat{f}(x)$, e.g. non-linear, can these predict better for the same test data set?

Polynomial Regression

A. Giovanidis 2022

The linear regression model assumes a linear relationship between the response and the input (predictors).

☞ But! the true relationship may be non-linear.

✎ Extend the linear model, using polynomial regression.

For 1-D input x we write an ℓ-polynomial model:

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \ldots + \beta_\ell x^\ell + \epsilon = f(x) + \epsilon. \tag{2}$$

But it is still a linear model for the parameters!

☞ If we regard $x_1 := x$, $x_2 := x^2, \ldots, x_\ell := x^\ell$ it is just a multiple linear regression.

$\rightarrow$ Use standard linear regression software.
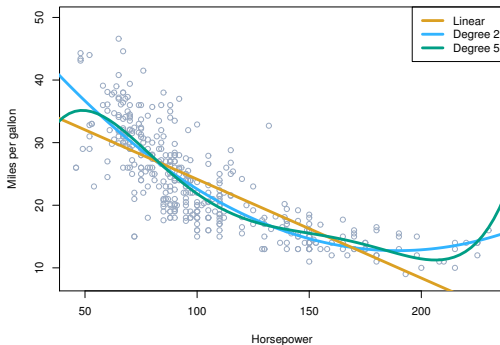
Polynomial Fit

A. Giovanidis 2022

Figure: Polynomial vs Linear fit.[1]

---
[1] Source [B.1]

Flexibility VS Interpretability

A. Giovanidis 2022

Higher order polynomial models offer more flexibility.
(see Question 2)

☞ In the most extreme case we can use a model whose curve passes through every point of the train data set D_n. We can propose a polynomial fit with $\ell = card(D_n)$.
Is this a good predictive model?

☞ At the other extreme we can use a very restrictive model (simple linear regression), with $\ell = 1$.
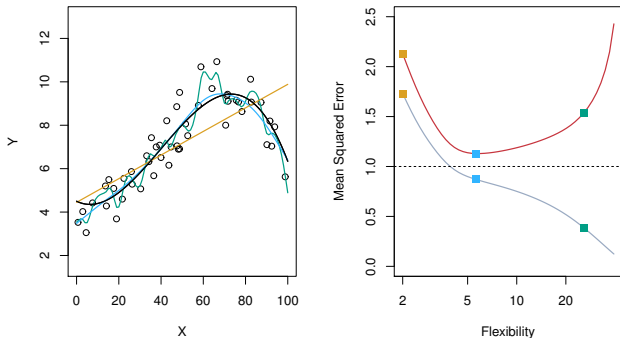Maybe this simple model is better?

Model Fitting I

A. Giovanidis 2022

Figure: Example 1 (black curve is the real one, noise added).[2]

---

[2]Source [B.1]

Model Fitting II

A. Giovanidis 2022

Figure: Example 2 (black curve is the real one, noise added).[3]

---

[3]Source [B.1]

Model Fitting III

A. Giovanidis 2022

Figure: Example 3 (black curve is the real one, noise added).[4]

---

[4]Source [B.1]

Optimal mode selection

Answer to Question 1:

☞ The MSE_{test} is always higher than noise (irreducible error).

Answer to Question 2:

The more flexibility (higher poly-degree ℓ), the lower the MSE_{train}.

But! The MSE_{test} always has a **U** shape (fundamental property) with respect to degree (x-axis). The optimal mode is the one that minimizes the MSE_{test}: trade-off between flexibility vs interpretability.

☞ We call this the Variance VS Bias trade-off.

Overfitting / Underfitting

A. Giovanidis 2022

Overfitting: Small MSE_{train} but large MSE_{test}. The statistical learning model picks patterns that are caused by randomness rather than the true properties of the unknown $f(x)$.

☞ needs lower flexibility!

Underfitting: Large MSE_{train} and large MSE_{test}. The learning model is too rigid to accurately describe the unknown $f(x)$.

☞ needs higher flexibility!

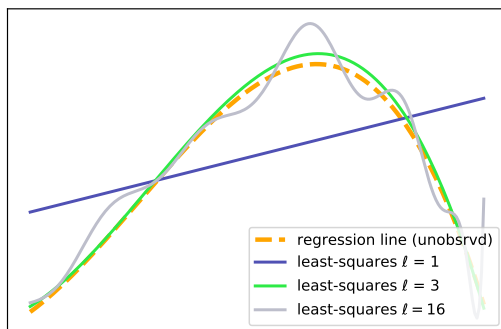# Overfitting / Underfitting Example

A. Giovanidis 2022

Figure: Polynomial Over/Under- fitting.

$MSE(Train) = (46.41, 25.66, 24.67)$, min train MSE for $\ell = 16$.
$MSE(Test) = (59.30, 42.94, 45.70)$, min test MSE for $\ell = 3$.

Numerical Example - Polynomial fit

A. Giovanidis 2022

Consider the numerical example with
$D_n = \{(1, 3),\ (2, 4),\ (3, 8),\ (4, 14)\}$.

- ► With the first three set elements do linear regression.

- ► Use the forth element to derive the $MSE_{test-linear}$.

- ► With the first three set elements do quadratic regression.

- ► Use the forth element to derive the $MSE_{test-quadratic}$.

<div align="center">Which method is best?</div>

Numerical Example - Polynomial fit cont'd ^{A. Giovanidis 2022}

We use the set $D_3 = \{(1,3),\ (2,4),\ (3,8))\}$ for the linear regression, and we get:

- $\hat{y} = 2.50x + 0$, with an $MSE_{test-linear} = 16$.

for the quadratic regression $\beta = (\mathbf{XX}^T)^{-1} \mathbf{X}^T \mathbf{y}$, and we get:

- $\hat{y} = 1.50x^2 - 3.5x + 5$, with an $MSE_{test-quadratic} = 1$.

$MSE_{test-quadratic} < MSE_{test-linear}$ ☞ the quadratic fit is better!

Train - Validation -Test sets

A. Giovanidis 2022

In all the above, we assumed that two disjoint data-sets are available:

- ▶ a train data-set D_n,

- ▶ a test data-set D_m^{test},

where $D_n \cap D_m^{test} = \emptyset$.

However, usually we only have one data-set available D_n, to both train and test the machine learning algorithm.

What should we do? a. Cross-validation, or b. Bootstrapping !

The data subsets

▶ training set: The sample of data used to fit the model. (e.g. weights and biases in the linear model or neural network)

▶ validation set: The sample of data used to provide an unbiased evaluation of a model fit on the training dataset while fine-tuning model hyper-parameters. (frequent evaluations)

▶ test set: The sample of data used to provide an unbiased evaluation of a final model fit on the training dataset. (only used once a model is completely trained)

Sometimes we can use the validation set to test our model (as a test set) but is not good practice.

Dataset Split

A. Giovanidis 2022

A visualization of the splits

Figure: Source: https://towardsdatascience.com/train-validation-and-test-sets-72cb40cba9e7

Validation set approach

☞ Consider that a test set is provided independently of observed data.

Naive approach. Split the observation data set D_n in two:

▶ a train set

▶ a validation set

e.g. 80% of the elements of D_n belong to the train and the rest 20% to the validation set.

▶ Use the train set to fit the model.

▶ Use the validation set to evaluate performance, e.g. MSE_{test}.

Validation set Example

A. Giovanidis 2022

Figure: *n* observations randomly split into a Train and Validation set.[5]

---

[5] Source [B.1]

Drawbacks

A. Giovanidis 2022

The method has two main problems:

1. The test error rate depends on the data split
   $\rightarrow MSE_{test}$ can be highly variable !

2. Not all available n data are used for training
   $\rightarrow$ worse performance with less observations, and MSE_{test} is overestimated !

# Validation set Example

A. Giovanidis 2022

Figure: Variability of the MSE_{test} depending on the data split.[6]

---

[6]Source [B.1]

LOOCV

Leave-one-out-Cross-Validation. Split the data set D_n again in two:

- a validation set of a single observation (x_1, y_1)
- a train set of the rest $n - 1$ observations .

A prediction $\hat{y}_1$ is made only for the excluded observation using x_1.

$$MSE_{test} = MSE_1 = (y_1 - \hat{y}_1)^2.$$

- Problem: The evaluation is based on a single observation $\rightarrow$ highly variable.

LOOCV x n

A. Giovanidis 2022

Figure: Solution: repeat *n* times, for *n* different splits ![7]

---

[7] Source [B.1]

LOOCV x n

1	2 3 … n-1 n	$\rightarrow MSE_1$
2	1 3 … n-1 n	$\rightarrow MSE_2$
3	1 2 … n-1 n	$\rightarrow MSE_3$
…	…	…
n	1 2 3 … n-1	$\rightarrow MSE_n$

The LOOCV estimate for the test MSE is the average of these n test error-estimates

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^{n} MSE_i.$$

☞ No randomness in the result $+$ uses all observations !

k-fold CV

For large n LOOCV is time-consuming to calculate all MSE_i.

☞Better use k-fold Cross-Validation:

- The data-set D_n is split into $1 \leq k \leq n$ folds.
- The 1st fold is treated as validation set and the rest $k - 1$ for training $\rightarrow MSE_1$ is caclulated.
- Repeat k times by choosing a different fold for validation set each time $\rightarrow MSE_k$ is calculated.

The k-fold estimate is computed by averaging

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^{k} MSE_i.$$

k-fold x k

A. Giovanidis 2022

Figure: Repeat k times, for k different splits ![8]

---

[8] Source [B.1]

k-fold CV

Computational advantages over LOOCV:

- LOOCV is a special case of k-fold CV, for $k = n$.

- In practice $k = 5$ or $k = 10$.

- k-fold CV with $k < n$ can have lower variance in the *MSE*, than LOOCV.

LOOCV vs k-fold CV

A. Giovanidis 2022

Figure: Mode selection with two CV methods[9].

The 10-fold was run 10 times, each with a different data split.

---

[9] Source [B.1]

The Bootstrap

A. Giovanidis 2022

A powerful tool that can quantify the uncertainty associated with a given learning method.

e.g. it can estimate the standard error (SE) of $\hat{\beta}_k$, or MSE_{test},...

Main idea:

Given an original data-set $Z = D_n$, create $B > 1$ new datasets of size n:

☞ each new dataset Z^{*b} results from uniform sampling with replacement of the set Z.

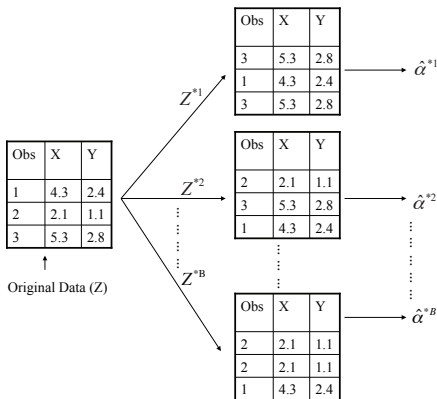☞ for each Z^{*b} calculate the unknown $\hat{\alpha}^{*b}$.

Resampled Data Sets

A. Giovanidis 2022

Figure: Example of Z^{*b} sets, $b = 1, 2, \ldots, B$.[10].

---

[10] Source [B.1]

Bootstrap estimates

▶ The Average

$$Av_B(\hat{\alpha}) = \bar{\alpha} = \frac{1}{B} \sum_{b=1}^{B} \hat{\alpha}^{*b}$$

▶ The Standard Error

$$SE_B(\hat{\alpha}) = \sqrt{\frac{1}{(B-1)} \sum_{b=1}^{B} (\hat{\alpha}^{*b} - \bar{\alpha})}$$

Resampled stats

A. Giovanidis 2022

Figure: Real vs Bootstrapped statistics.[11].

---

[11]Source [B.1]

Numerical Example - Resampling

A. Giovanidis 2022

Consider the numerical example with $D_n = \{(1,3),\ (2,4),\ (3,8),\ (4,9)\}$.

Use this set and apply :

- ▶ Simple Validation.
- ▶ LOOCV.
- ▶ 2-fold CV.
- ▶ Bootstrapping.

Numerical Example - Simple Validation

A. Giovanidis 2022

Suppose first that we split the D_n in half, into a

▶ Train Set $\{(1,3),\ (2,4)\}$.

▶ Validation Set $\{(3,8),(4,9)\}$.

Then $\hat{y} = 1 \cdot \hat{x} + 2$.

▶ $MSE_{train} = 0$.

▶ $MSE_{test} = \frac{(5-8)^2+(6-9)^2}{2} = 9$.

Numerical Example - LOOCV

A. Giovanidis 2022

We apply 4 times LOOCV

	Train	Validate
Fold 1	$\{(2,4),\ (3,8),\ (4,9)\}$	$(1,3)$
Fold 2	$\{(1,3),\ (3,8),\ (4,9)\}$	$(2,4)$
Fold 3	$\{(1,3),\ (2,4),\ (4,9)\}$	$(3,8)$
Fold 4	$\{(1,3),\ (2,4),\ (3,8)\}$	$(4,9)$

We get

Fold 1	$\hat{y} = 2.50x - 0.50$	$MSE_{test,1} = 1.$
Fold 2	$\hat{y} = 2.07x + 1.14$	$MSE_{test,2} = 1.638$
Fold 3	$\hat{y} = 2.07x + 0.50$	$MSE_{test,3} = 1.664$
Fold 4	$\hat{y} = 2.50x + 0.00$	$MSE_{test,4} = 1.$

$Av(MSE_{test}) = \frac{1+1.638+1.664+1}{4} = 1.3255$

40 / 46

Numerical Example - 2Folds

A. Giovanidis 2022

We apply 2-Folds (with shuffle)

	Train	Validate
Fold 1	$\{(3,8),\ (2,4)\}$	$\{(1,3),\ (4,9)\}$
Fold 2	$\{(1,3),\ (4,9)\}$	$\{(3,8),\ (2,4)\}$

We get

Fold 1	$\hat{y} = 4.00x - 4.00$	$MSE_{test,1} = 9$
Fold 2	$\hat{y} = 2.00x + 1.00$	$MSE_{test,2} = 1$

$Av(MSE_{test}) = \frac{9+1}{2} = 5$.

Numerical Example - Bootstrapping

A. Giovanidis 2022

Consider the numerical example with
$D_n = Z_1 = \{(1,3),\ (2,4),\ (3,8),(4,9)\ \}$.

Apply the bootstrapping technique for 3 more sample sets :

- $Z_2^* = \{(2,4),\ (2,4),\ (4,9),(4,9)\ \}$.
- $Z_3^* = \{(2,4),\ (1,3),\ (3,8),(4,9)\ \}$.
- $Z_4^* = \{(2,4),\ (4,9),\ (3,8),(3,8)\ \}$.

Numerical Example - Bootstrapping cont'd IA. Giovanidis 2022

We get for each of the four sets,

1. $(\hat{\beta}_0, \ \hat{\beta}_1) = (0.5, \ 2.2)$.
2. $(\hat{\beta}_0, \ \hat{\beta}_1) = (-1.0, \ 2.5)$.
3. $(\hat{\beta}_0, \ \hat{\beta}_1) = (0.5, \ 2.2)$.
4. $(\hat{\beta}_0, \ \hat{\beta}_1) = (-0.25, \ 2.5)$.

Numerical Example - Bootstrapping cont'd II

A. Giovanidis 2022

Figure: Bootstrap Regression Lines.

Numerical Example - Bootstrapping cont'd III

A. Giovanidis 2022

We have obtained the following set of estimates for the intercept and the slope

- $\hat{\beta}_0 = [0.5, \ -1.0, \ 0.5, \ -0.25 \ ]$.
- $\hat{\beta}_1 = [2.2, \ 2.5, \ 2.2, \ 2.5 \ ]$.

Then

- $\bar{\beta}_0 = Av_4(\hat{\beta}_0) = -0.0625$ and $SE_4(\hat{\beta}_0) = 0.3590$.
- $\bar{\beta}_1 = Av_4(\hat{\beta}_1) = 2.35$ and $SE_4(\hat{\beta}_1) = 0.0867$.

95% confidence intervals:

- $\beta_0 \in [-0.7806, \ +0.6556]$.
- $\beta_1 \in [+2.1768, \ +2.5232]$.

A. Giovanidis 2022

END